

Pydiogment: A Python package for audio augmentation

Ayoub Malek¹ and Hasna Marwa Malek²

DOI: [00.00000/joss.00000](https://doi.org/00.00000/joss.00000)

¹ Technical University of Munich ² Grenoble Institute of Technology

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Submitted: 25 February 2020

Published: 25 February 2020

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC-BY](#)).

Summary

Audio data augmentation is a key step in training Machine Learning (ML) models to solve audio classification tasks. It is applied to increase the quality and size of the labeled training data-set, in order to improve the recognition accuracy. Data augmentation is simply a deformation technique, that helps stretch the data, and increase its size for a better training. Unlike image augmentation, audio augmentation is still limitly explored by research and most deformation strategies manipulate the computed spectrograms rather than the raw audio. With the exception of few libraries constrained to work with Pytorch (Paszke et al., 2019), most existing tools in this context either act on spectrograms such as Google's SpecAugment (Park et al., 2019), or are developed for music data augmentation like muda (McFee, Humphrey, & Bello, 2015). This paper describes version 0.1.0 of **Pydiogment**: a Python package for audio augmentation based on the Scipy (???) and ffmpeg (FFmpeg Developers, 2019) libraries. **Pydiogment** implements various augmentation techniques that can be used to improve the accuracy of various recognition tasks (speaker recognition, spoken emotions recognition, speech recognition etc.) and avoid over-fitting when training models. The paper provides a brief overview of the library's functionality, along with a small emotions recognition experiment displaying the utility of the library.

Implementation and theory

Pydiogment includes 3 general categories of deformations / augmentations:

Amplitude based augmentations (auga.py)

- **Apply Gain:** This deformation can be described as an amplification of the signal and the noise by applying a given gain (in dB) to the input signal. Note that excessive gain application can result in clipping (Self et al., 2009).
- **Add Fade:** adds fade-in and fade-out effects to the original signal. This is done by multiplying a hamming window with the original signal $y[n] = x[n] * w[n]$ where $x[n]$ is the original signal, $y[n]$ is the augmented signal and $w[n]$ is the computed hamming window (Poularikas, 1999).
- **Normalize:** Normalization refers to the practice of applying a uniform amount of gain across a signal, where signal-to-noise ratio and general dynamics levels remain unchanged (Shelvock, 2012). The normalization can be applied using the peak normalization method $y[n] = \frac{x[n]}{\max(x[n])}$ or the Root Mean Square (RMS) approach

$y[n] = \sqrt{\frac{N \cdot 10^{\frac{r}{20}}}{\sum_{n=0}^{N-1} x^2[n]}} \cdot x[n]$ where $x[n]$ is the original signal, $y[n]$ is the augmented signal, N is the length of $x[n]$ and r is the input RMS level in dB.

- **Add Noise:** Additive White Gaussian Noise (AWGN) is added to the input signal based on a given signal-to-noise ratio (SNR) in dB :

$$y[n] = x[n] + \sqrt{\frac{P_x}{P_{awgn}} \cdot 10^{-\left(\frac{SNR_{db}}{10}\right)}} \cdot awgn[n] \quad (1)$$

with $x[n]$ is the original signal, $y[n]$ is the augmented noisy signal, $awgn[n]$ is a random Gaussian white noise signal with standard deviation = 1 & mean = 0, P_x & P_{awgn} are respectively the signal power and noise power (Balakrishnan & Verghese, 2012).

Frequency based augmentation (augf.py)

- **Change tone:** a tone or pitch are sound's properties that allow ordering on a frequency-related scale. This is usually characterized by a fundamental frequency that can be adjusted without changing the tempo to provide a deformed audio, this is also know as pitch shifting (Bernsee, 2005).
- **Convolve:** This is also called reverberating the audio and it consists of convolving the original signal with a given Room Impulse Response (RIR) to simulate an audio captured using far-field microphones in a different setup/channel $y[n] = x[n] * rir[n]$ where $x[n]$ is the original signal, $y[n]$ is the augmented signal and $rir[n]$ is the room impulse response (Raju, Panchapagesan, Liu, Mandal, & Strom, 2018).
- **Apply Filter:** apply various types of Scipy based Butterworth filters (low-pass, high-pass or band-pass).

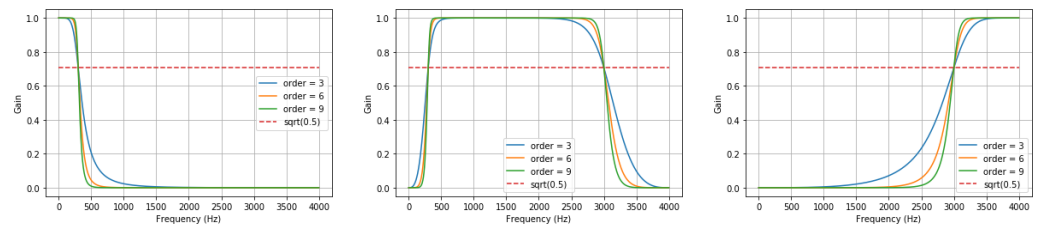


Figure 1: Frequency-Gain graphs of the implemented Butterworth filers (left: low-pass, middle: band-pass, right: high-pass) with a low cutoff frequency of 300 Hz and a high cutoff frequency of 3000 Hz .

Time based augmentation (augt.py)

- **Stretch Time:** also known as time compression/expansion, it is reciprocal to pitch shifting. Essentially the audio is slowed-down or accelerated based on a given coefficient (Bernsee, 2005).
- **Shift Time:** this includes shifting the signal in a certain time direction to create augmented signals with different chronological orders, and also reversing the whole signal.
- **Crop:** generates a randomly cropped audio based on the original signal and minimum signal/audio length.
- **Remove Silence:** filters out silent frames from the input signal using FFmpeg silenceremove filter (FFmpeg Developers, 2019).

- **Resample:** with the help of Scipy, the input signal is resampled given an input sampling rate with respect to the Nyquist–Shannon sampling theorem.

The aforementioned augmentation strategies can be combined to generate various sub-strategies. However, it is very crucial to maintain the semantic validity when augmenting the data. *For example:* one cannot change tones when doing voice based gender classification and still expect tone to be a separating feature of the predicted classes.

Experiment & Results

To prove the utility of **Pydiogment**, we display its effect on a spoken emotions recognition task. We use the **Emo-DB** data-set (Burkhardt, F., Paeschke, A., Rolfes, M., A., Walter F. Sendlmeier & Weiss, B., 2005) as a starting point, which is a small German audio data-set simulating 7 different emotions (neutral, sadness, anger, boredom, fear, happiness, disgust). We choose the Mel-Frequency Cepstral Coefficients (MFCCs) (Milner & Shao, 2006) as the characterizing low-level audio features due to previous proved success on similar problems (Dahake, Shaw, & Malathi, 2016; Kandali, Routray, & Basu, 2008; Krishna Kishore & Krishna Satish, 2013; Sreeram, Geyasruti, Narayanan, & M, 2015). The features are extracted using the `python_speech_features` library (Lyons et al., 2020). In a first phase and using the `scikit-learn` library (Pedregosa et al., 2011), we apply various recognition algorithms on the original data such as K-Nearest Neighbors (KNN), random forests, decision trees, Quadratic Discriminant Analysis (QDA), Support Vector Machines (SVM) etc. In a second phase, we augment the data using **Pydiogment** by applying the following techniques:

- slow down samples using a coefficient of 0.8.
- speed up samples coefficient of 1.2.
- randomly crop samples with a minimum length of 1 second.
- add noise with an SNR = 10
- add a fade in and fade out effect.
- apply gain of -100 dB.
- apply gain of -50 dB.
- convolve with noise file using a level = $10^{*-2.75}$.
- shift time with one second (1 sec) to the right (direction = right)
- shift time with one second (1 sec) to the left (direction = left)
- change tone with tone coefficient equal to 0.9.
- change tone with tone coefficient equal to 1.1.

Then we re-run the same recognition algorithms on the augmented and original data. The following is a comparison of the results:

Machine learning Algorithm	Accuracy (no augmentation)	Accuracy (with augmentation)
AdaBoost	0.309	0.513
Decision Tree	0.454	0.764
Extra Trees Classifier	0.588	0.916
Gaussian Process	0.247	0.700
KNN	0.629	0.867
Linear SVM	0.608	0.693
MLP Classifier	0.608	0.811
Naive Bayes	0.577	0.610
QDA	0.608	0.764
Random Forest	0.567	0.929

Table 1: Accuracy comparison of results with and without data augmentation.

Conclusion

This paper introduced `Pydiogment`, a Python package for audio data augmentation, with diverse audio deformation strategies. These strategies aim to improve the accuracy of audio based recognition system by scaling the training data-set and increasing its quality/diversity. The utility of `Pydiogment` was proved by showing its effects when used in a spoken emotions recognition task. In the stated experiment, the augmentation using `Pydiogment` improved the accuracy up to 50%.

References

- Balakrishnan, H., & Verghese, G. (2012). 6.02 Introduction to EECS II. *Digital Communication Systems*. Massachusetts Institute of Technology: MIT OpenCourseWare, <https://ocw.mit.edu>. License: [Creative Commons BY-NC-SA](#).
- Bernsee, S. M. (2005). Time stretching and pitch shifting of audio signals an overview. Retrieved from <http://blogs.zynaptiq.com/bernsee/time-pitch-overview/>
- Burkhardt, F., Paeschke, A., Rolfes, M., A., Walter F. Sendlmeier & Weiss, B. (2005). A database of german emotional speech. *INTERSPEECH*.
- Dahake, P. P., Shaw, K., & Malathi, P. (2016). Speaker dependent speech emotion recognition using mfcc and support vector machine. In *2016 international conference on automatic control and dynamic optimization techniques (icadot)* (pp. 1080–1084). doi:[10.1109/ICACDOT.2016.7877753](https://doi.org/10.1109/ICACDOT.2016.7877753)
- FFmpeg Developers. (2019). FFmpeg tool (version 4.2.2). Retrieved from <http://ffmpeg.org/>
- Kandali, A. B., Routray, A., & Basu, T. K. (2008). Emotion recognition from assamese speeches using mfcc features and gmm classifier. In *TENCON 2008 - 2008 ieee region 10 conference* (pp. 1–5). doi:[10.1109/TENCON.2008.4766487](https://doi.org/10.1109/TENCON.2008.4766487)
- Krishna Kishore, K. V., & Krishna Satish, P. (2013). Emotion recognition in speech using mfcc and wavelet features. In *2013 3rd ieee international advance computing conference (iacc)* (pp. 842–847). doi:[10.1109/IAdCC.2013.6514336](https://doi.org/10.1109/IAdCC.2013.6514336)
- Lyons, J., Wang, D. Y.-B., Gianluca, Shteingart, H., Mavrinac, E., Gaurkar, Y., Watchara-wisetkul, W., et al. (2020). *Jameslyons/python_speech_features: Release v0.6.1*. Zenodo. doi:[10.5281/zenodo.3607820](https://doi.org/10.5281/zenodo.3607820)
- McFee, B., Humphrey, E., & Bello, J. (2015). A software framework for musical data augmentation. In *16th international society for music information retrieval conference, ISMIR*. Retrieved from <https://github.com/bmcfee/muda>
- Milner, B., & Shao, X. (2006). Clean speech reconstruction from mfcc vectors and fundamental frequency using an integrated front-end. *Speech Communication*, 48(6), 697–715. doi:[10.1016/j.specom.2005.10.004](https://doi.org/10.1016/j.specom.2005.10.004)
- Park, D. S., Chan, W., Zhang, Y., Chiu, C.-C., Zoph, B., Cubuk, E. D., & Le, Q. V. (2019). SpecAugment: A simple augmentation method for automatic speech recognition. In *INTERSPEECH*. doi:[10.21437/interspeech.2019-2680](https://doi.org/10.21437/interspeech.2019-2680)
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., et al. (2019). PyTorch: An imperative style, high-performance deep learning library. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, & R. Garnett (Eds.), *Advances in neural information processing systems 32* (pp.

8024–8035). Curran Associates, Inc. Retrieved from <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., et al. (2011). Scikit-learn: Machine learning in python. *J. Mach. Learn. Res.*, 12, 2825–2830.

Poularikas, A. D. (1999). *The handbook of formulas and tables for signal processing* (1st ed., Vol. 1, pp. 137–138). Boca Raton, USA: CRC Press. doi:[10.1201/9781315219707](https://doi.org/10.1201/9781315219707)

Raju, A., Panchapagesan, S., Liu, X., Mandal, A., & Strom, N. (2018). Data augmentation for robust keyword spotting under playback interference. *ArXiv, abs/1808.00563*.

Self, D., Duncan, B., Sinclair, I., Brice, R., Hood, J. L., Singmin, A., Davis, D., et al. (2009). *Audio engineering: Know it all* (Vol. 1, pp. 277–278). Oxford, UK: Elsevier Inc.

Shelvock, M. (2012). *Audio Mastering as Musical Practice* (Master's thesis). The University of Western Ontario: The School of Graduate; Postdoctoral Studies, London, Ontario, Canada. Retrieved from <https://ir.lib.uwo.ca/etd/530>

Sreeram, L., Geyasruti, D., Narayanan, R., & M, S. (2015). Emotion detection using mfcc and cepstrum features. *Procedia Computer Science*, 70, 29–35. doi:[10.1016/j.procs.2015.10.020](https://doi.org/10.1016/j.procs.2015.10.020)